
Columnar Indexes for Row Reconstruction

E0261

Jayant Haritsa
Indian Institute of Science



Motivation

- How can we make Tuple Reconstruction easier?
 - if columns are sorted in the **same** order
 - Eg: Discount BAT is sorted on Discount, and Day BAT is reorganized to have matching tuple-ID sequence

Select Day, Discount

From Sales

Where Discount $\leq$ 0.15

- However, the sorted tuple-ordering cannot always be preserved
 - During query processing, many operators (joins, group by, order by, etc.) are not **tuple order-preserving**

Stratos Idreos




*Won the 2011 ACM SIGMOD Jim Gray
Doctoral Dissertation Award for his thesis
'Database Cracking: Towards Auto-tuning
Database Kernels'.*

Database Cracking

- Self-tuning database kernel
- Every query is treated as an advice on how data should be stored
 - Physical reorganization happens per column based on selection predicates

Select R.A
From R
Where R.A <10



Cracking: Example [CIDR 2007]






Note: Data portion not used is not tuned

Q1:
select R.A
from R
where R.A > 10
and R.A < 14

Q2:
select R.A
from R
where R.A > 7
and R.A ≤ 16

column A

13
16
4
9
2
12
7
1
19
3
14
11
8
6



result

Selection-Based Cracking

- The first time an attribute A is required by a query, a copy of column A is created, called the **cracker column** C_A of A
- Each selection operator on A **triggers** a range-based physical reorganization of C_A
 - *Implemented similar to **bubble-sort** across the beginning and end segments of the cracked portion* [Algo 2 of CIDR07 paper]
- Each cracker column, has a **cracker index** (AVL-tree) to maintain partitioning (piece) information
- Future queries benefit from the physically clustered data and do not need to access the whole column

The `crackers.select` Operator

`crackers.select(A, v1, v2)`

- First, it creates C_A if it does not exist
- It searches the index of C_A for the area where $v1$ and $v2$ fall
- If the bounds do not exist, i.e., no query used them in the past, then C_A is physically reorganized to cluster all qualifying tuples into a contiguous area
- Output:
 - A list of keys/positions

But, how to handle queries such as
`Select R.B From R where R.A <10`

Cracker Map

- For Query Type
 - Access attribute B based on attribute A of same relation R
- A cracker map M_{AB} is defined as a two-column table over two attributes A and B of relation R
 - Values of A are stored in the left column, called **head**
 - Values of B are stored in the right column, called **tail**
- Values of A and B in the *same* position of M_{AB} belong to the *same* tuple of R

Maps are Created on Demand Only

- When a query q needs access to attribute B based on a restriction on attribute A and M_{AB} does not exist,
 - then q will create M_{AB} by performing a scan over base columns A and B
- Else, if M_{AB} already exists,
 - All tuples with values of A that qualify the restriction are in a contiguous area in M_{AB}
 - Realized by splitting a piece of M_{AB} into two or three new pieces
- For each cracker map M_{AB} , there is a cracker index (AVL-tree) that maintains information about how A values are distributed over M_{AB}

Example

Initial state

A	B
12	b1
3	b2
5	b3
9	b4
15	b5
22	b6
7	b7
26	b8
4	b9
2	b10
24	b11
11	b12
16	b13

select B from R where $10 < A < 15$

Cracker index		M_{AB}	
Piece 1	Position 1 value ≤ 10	4	b9
		3	b2
		5	b3
		9	b4
		2	b10
Piece 2	Position 7 value > 10	7	b7
		12	b1
		11	b12
		15	b5
Piece 3	Position 9 value ≥ 15	22	b6
		24	b11
		26	b8
		16	b13

select B from R where $5 \leq A < 17$

Cracker index		M_{AB}	
Piece 1	Position 1 value < 5	4	b9
		3	b2
		2	b10
Piece 2	Position 4 value ≥ 5	9	b4
		5	b3
		7	b7
Piece 3	Position 7 value > 10	12	b1
		11	b12
Piece 4	Position 9 value ≥ 15	15	b5
		16	b13
		24	b11
Piece 5	Position 11 value ≥ 17	26	b8
		22	b6

Sideways.select(A , $v1$, $v2$, B) Operator

Returns tuples of attribute B of relation R based on a predicate on attribute A of R as follows:

- (1) If there is no cracker map M_{AB} , then create one
- (2) Search the index of M_{AB} to find the contiguous area w of the pieces related to the restriction σ on A

If σ does not match existing piece boundaries,



- (3) Physically reorganize w to move false hits out of the contiguous area of qualifying tuples
- (4) Update the cracker index of M_{AB} accordingly
- (5) Return a non-materialized view of the tail of w

Multiple Projection Queries

- How to handle the following query?
Select B, C
From R
Where $A < 10$
- For this query, we need 2 maps M_{AB} and M_{AC}
- In general, a single-selection query q that projects n attributes requires n maps, one for each attribute to be projected
- All maps that have been created using A as head are collected in the **map set** S_A


Initial state

A	B	C
7	b1	c1
4	b2	c2
1	b3	c3
2	b4	c4
8	b5	c5
3	b6	c6
6	b7	c7



The Problem


Naïve use of the sideways.select operator may lead to non-aligned cracker maps



Solution: Adaptive Alignment

- Extend the sideways.select operator with an **alignment step** to keep the maps aligned
- The Basic Idea
 - is to apply **all** physical reorganizations, due to selections on an attribute **A** , in the *same order* to all maps in the **map set S_A**





Cracker Tape

- For each map set S_A , introduce a cracker tape T_A
 - T_A logs (in order of their occurrence) all selections on attribute A that trigger cracking of any map in S_A
 - Each map M_{Ax} is equipped with a *cursor* pointing to the entry in T_A that represents the last crack on M_{Ax}
- Given a tape T_A , a map M_{Ax} is aligned (synchronized) by successively forwarding its cursor towards the end of M_{Ax}
- And incrementally cracking M_{Ax} according to all selections it passes on its way
- All maps whose cursors point to the same position in T_A , are physically aligned

The Extended sideways.select Operator: Summary

- (1) If there is no T_A , then create an empty one.
- (2) If there is no cracker map M_{AB} , then create one.
- (3) Align M_{AB} using T_A .
- (4) Search the index of M_{AB} to find the contiguous area w of the pieces related to the restriction σ on A .
If σ does not match existing piece boundaries,
 - (5) Physically reorganize w to move false hits out of the contiguous area of qualifying tuples.
 - (6) Update the cracker index of M_{AB} accordingly.
 - (7) Append predicate $v_1 < A < v_2$ to T_A .
- (8) Return a non-materialized view of the tail of w .

Multiple-Select Queries

Query: **Select** D **From** R
 Where $3 < A < 10$ and
 $4 < B < 8$ and $1 < C < 7$

First-cut solution

- Create different maps M_{AD}, M_{BD}, M_{CD}
- Does this work?
- No, because of distinct sorting orders!

Solution

Query: **Select D From R**
Where $3 < A < 10$ and
 $4 < B < 8$ and $1 < C < 7$


- Choose a selection predicate, say **A** (tuple ordering is possible only on one predicate)
- Create Maps M_{AB} , M_{AC} and M_{AD}
- Since **B** is also a selection predicate in the Map M_{AB} , create a bit vector for **B** which identifies the tuples satisfying predicate **B**
- Same is repeated for M_{AC}
- Use ANDing of the two bit vectors with Map M_{AD} to get the desired result

Query


Initial state

select D from R
where $3 < A < 10$ and
 $4 < B < 8$ and $1 < C < 7$


A	B	C	D
12	9	3	9
3	2	6	4
5	6	2	2
9	10	1	10
8	7	6	12
22	11	9	19
7	16	12	3
26	2	2	6
4	5	11	5
2	8	17	8
7	3	3	1




select_create_bv(A,3,10,B,4,8)



select_refine_bv(A,3,10,C,1,7,bv)



reconstruct(A,3,10,D,bv)



16	2
2	3
2	9

Example

Map Set Choice: Self-organizing Histograms

- We had chosen predicate A randomly to be head for all the Maps. Is there a better choice of the predicate?
- Yes! for a query q , a set S_A is chosen such that the restriction on A is the most selective in q
 - yielding a minimal-size bit vector
- The most selective restriction can be found using the cracker indices
 - choose the most aligned Map M_{AX} (for accurate statistics)
 - Estimate cardinality interval of the predicate using index boundaries

Complex Queries

- Apart from tuple reconstruction, other operators do not depend on **tuple insertion order**
 - Not affected by cracking physical reorganization
 - Joins, aggregations, groupings, etc.
 - Therefore, use standard column store operators
- Potentially many operators can exploit the clustering information in the maps
 - MAX operator can consider only the last piece of a map



Experimental Analysis

- Compare the implementation of selection and sideways cracking on top of MonetDB
- Against the non-cracking version of MonetDB
- And against MonetDB on presorted data
- Results
 - Sideways cracking achieves similar performance to presorted data
 - But does not have the heavy initial cost and the restrictions on updates and workload prediction



Partial Sideways Cracking

- Consider storage restriction
- Partial Maps
 - Maps are only partially materialized driven by the workload
 - A map consists of several **chunks**
 - Each chunk is a separate two-column table
 - Each chunk contains a given value range of the head attribute of this map
 - Each chunk is cracked separately



Partial Map Example




Figure 8: Using partial maps (U=Unfetched, F=Fetched, E=Empty, M=Materialized, C=ChunkID)

END COLUMNAR INDEXES

